



# Unit5


SJMurchite\_SYB\_Unit5 syllabus

Floating point and SIMD instructions

# Floating point Representation

- There are three kinds of floating point numbers in IA32 architectures.
- IEEE754 Single precision floating point format
- IEEE754 Double precision floating point format
- Double extended precision floating point number

Single



Double



Extended



# Floating point Representation

- IA32 architectures support IEEE754 number representation schemes of floating point numbers. There are three kinds of floating point numbers in IA32 architecture.
- **IEEE754 Single precision floating point format**
- These are 32 bit wide numbers and have three parts-a sign, a mantissa and an exponent.
- Sign is 1 bit wide and represents the sign of the number. Mantissa is 23 bit wide and represents the fractional part of the floating point number in binary. A leading one bit integer is not stored and is assumed as 1 in the normalized numbers.
- Exponents are stored in excess-127 representation.
- **IEEE754 Double precision floating point format**
- These are 64 bit wide number with sign, mantissa and exponent in excess 1023 representation.
- **Double extended precision floating point number**
- Double extended precision floating point numbers are 80 bit wide.
- This number format is used to perform computations internally in IA32 processors when instructions in x87 floating point instruction set are used.

## Floating point conversion in single precision

- Conversion of decimal number to IEEE single precision floating point format
- -24.75

Sign bit(S) = 1 (for negative number)

$$24.75 = 11000.11 * 2^0$$

$$1.100011 * 2^4$$

— Exponent is excess-127 in single precision format

$$\begin{aligned}\text{bias exponent} &= \text{actual exponent} + \text{bias value} \\ &= 127 + 4 = 131\end{aligned}$$

$$\text{Mantissa (M)} = (1.100011)$$



$$-24.75 = 0xC1C60000$$

$$+7.5 = ?$$

# Floating point conversion in double precision

- Conversion of decimal number to IEEE single precision floating point format
- -24.75

Sign bit(S) =1 (for negative number)

$$24.75 = 11000.11 * 2^0$$

$$1.100011 \cdot 2^4$$

— Exponent is excess-1023 in double precision format

$$\begin{aligned}\text{bias exponent} &= \text{actual exponent} + \text{bias value} \\ &= 1023 + 4 = 1027\end{aligned}$$

Mantissa (**M**)=(1.100011)

[illegible]

-24.75 = 0xC038 C000 0000 0000

$+7.5 = ?$

# First Floating point program

```
.section .data
```

```
value1:
```

```
.float -24.75
```

```
value2:
```

```
.double -24.75
```

```
.section .bss
```

```
.lcomm data,4
```

```
.lcomm data1,8
```

```
.section .text
```

```
.globl _start
```

```
_start : nop
```

```
fini
```

```
flds value1
```

```
fldl value2
```

```
fsts data
```

```
fstl data1
```

```
movl $1,%eax
```

```
movl $0,%ebx
```

```
int $0x80
```

## Single and double precision Number in IEEE


```
Terminal
student@student-OptiPlex-3020: ~
student@student-OptiPlex-3020:~$ as -gstabs -o float1.o float1.s
student@student-OptiPlex-3020:~$ ld -o float1 float1.o
student@student-OptiPlex-3020:~$ gdb -q float1
Reading symbols from float1...done.
(gdb) break*_start+1
Breakpoint 1 at 0x8048075: file float1.s, line 11.
(gdb) run
Starting program: /home/student/float1
Breakpoint 1, _start () at float1.s:11
11      flds value1
(gdb) x/fx &value1
0x8049093:      0xc1c60000
(gdb) s
12      fldl value2
(gdb) x/gfx &value2
0x8049097:      0xc038c00000000000
(gdb) s
13      fstl data
(gdb) s
14      movl $1,%eax
(gdb) x/gfx &data
0x80490a0 <data>:      0xc038c00000000000
(gdb)
```

# Double extended precision Number in IEEE

```
Terminal
student@student-OptiPlex-3020: ~
eax      0x0      0
ecx      0x0      0
edx      0x0      0
ebx      0x0      0
esp      0xbffff110      0xbffff110
ebp      0x0      0x0
esi      0x0      0
edi      0x0      0
eip      0x804807b      0x804807b <_start+7>
eflags   0x202      [ IF ]
cs       0x73      115
ss       0x7b      123
ds       0x7b      123
es       0x7b      123
fs       0x0      0
gs       0x0      0
st0      -24.75      (raw 0xc003c60000000000000000)
st1      0           (raw 0x0000000000000000000000)
st2      0           (raw 0x0000000000000000000000)
st3      0           (raw 0x0000000000000000000000)
st4      0           (raw 0x0000000000000000000000)
st5      0           (raw 0x0000000000000000000000)
st6      0           (raw 0x0000000000000000000000)
---Type <return> to continue, or q <return> to quit---
```



# Architecture of floating point processor



- **X87 general purpose Data registers in IA32 processor**

It include 8 general purpose data registers, each 80 bit wide and capable of storing a real number in double extended precision format.


- As data is loaded into the FPU stack, the stack top moves downward in the registers.

- When eight values have been loaded into the stack, all eight FPU data registers have been utilized.

- If a ninth value is loaded into the stack, the stack pointer wraps around to the first register and replaces the value in that register with the new value

# Architecture of floating point processor

FPU Register Stack



- When an integer ,BCD, single and double precision operand is loaded into one of these data registers, the operand is implicitly converted to double extended precision format.
- In addition to data registers, x87 also includes three 16 bit registers which are used in controlling the computations.
- These registers are the following
  - X87 control register
  - X87 status register
  - X87 tag register

# Top of stack

Stack register	Data register when TOP=5	Data register when TOP=2
St(0)	R5	R2
St(1)	R4	R1
St(2)	R3	R0
St(3)	R2	R7
St(4)	R1	R6
St(5)	R0	R5
St(6)	R7	R4
St(7)	R6	R3

The register stack used registers in a circular buffer.

1) When TOP=5, then st (0) refers to x87 data register R5.

# Floating point exceptions

- **Invalid operation exception (IE)**
  - It is raised when operands of the instruction contain invalid values such as NaN.
  - Example: division of 0 by 0 should lead to NaN
- **Divide by zero exception (ZE)**
  - Occurs when divisor contains a zero and dividend is a finite number other than 0
- **De-normalized operand exception (DE)**
  - It is raised when one of the operands of an instruction is in denormalized form.
  - De-normalized numbers can be used to represent numbers with magnitudes too small to normalize (i.e. below  $1.0 \times 2^{-126}$ )
- **Numeric overflow exception (OE)**
  - Occurred when rounding result would not fit into destination operand.
  - *When a double precision floating point number is to be stored in single precision floating point format, a numeric overflow is set to have occurred.*
- **Numeric underflow exception (UE)**
  - Occur when the instruction generate result whose magnitude is less than the smallest possible normalized number.
- **Precision (Inexact result) exception (PE)**
  - Division of (1.0/12.0) results in recurring infinite sequence of bits.
  - This number cannot be stored in any format without loss of precision.

# X87 control register


During computation ,several floating point error conditions such as divide by zero occur

In case of such errors , x87 FPU can be programmed to raise floating point exception

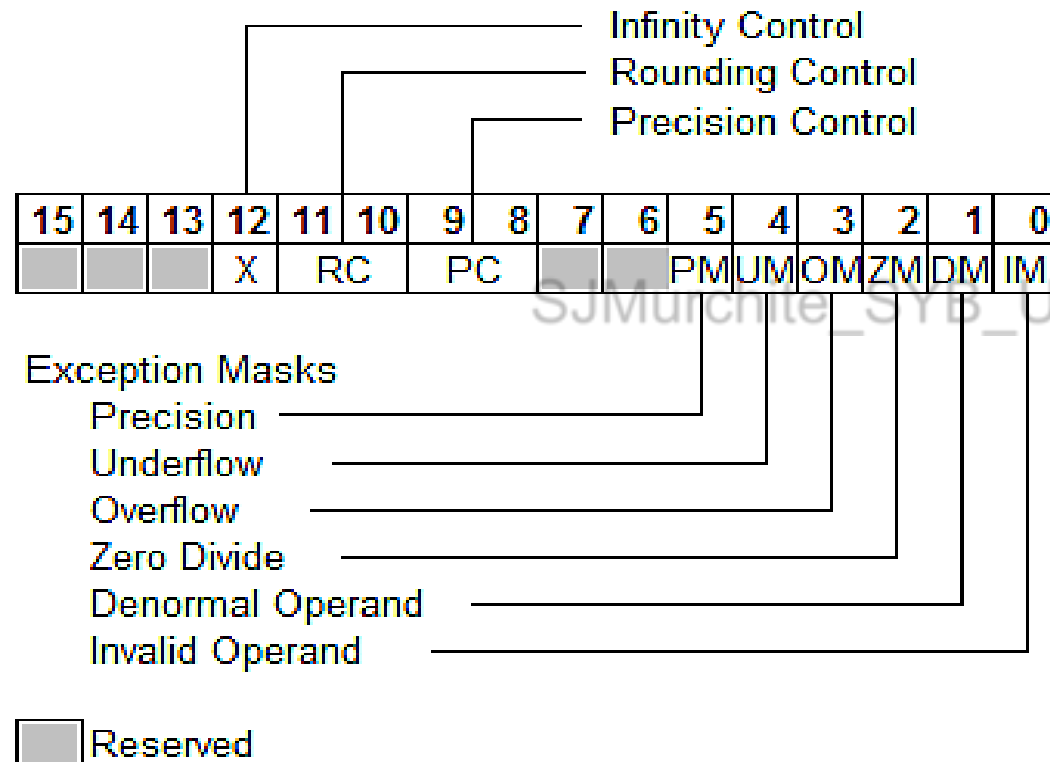
Bit 0 to bit 5 SJMurchite\_SYB\_Unit5 syllabus

0 indicate : unmask exception

1 indicate: Mask excpetion



# X87 control register



Used to handle the precision and rounding of operands

Infinity bit always set to 0

Precision control

00- single precision

01 unused

10-Double precision

11-Double extended precision

Rounding control

00-Round to nearest

01 Round down

10-Round up

11 Round towards zero(Truncate)

Exception Mask

0-Unmask

1-Mask

# Rounding Controls in IA 32 architectures

- +1.000100100011010011101 0011
  1. Rounding towards 0 +1.000100100011010011101
  2. Rounding towards  $+\infty$  (Rounding up) +1.000100100011010011110
- Rounding towards plus infinity control used 0011 would be removed but a 1 will be added to the LSB of remanding bits because the removed bit pattern is other than a zero.
- +1.000100100011010011101 1011
  3. Rounding towards int +1.0001001000110100011110
  4. Rounding towards  $-\infty$  (Rounding down) +1.0001001000110100011101

Note:1) Rounding towards 0 identical as rounding towards  $-\infty$

Note:1) Rounding towards int identical as rounding towards  $+\infty$

# X87 Status register

**1) TOP**-it is 3 bit field contain index to register assumed to be at the top of the stack.

**2) Four conditional code flags C0 to C3**

A floating point number can be compared in different ways with other floating point number.

The result of comparison is stored in either eflags or in floating point conditional code c0-c3

**3) SF**-The stack fault bit in status register indicates errors due to stack overflow or underflow conditions

Example: when one data on stack and addition operation is performed which takes two data from the stack, a stack fault occurs.

Stack overflow; when all 8 registers are in use and an attempt is made to load an data , stack overflow occurs.


**ES- Error status**

1- Error occurred    0- No error

**B bit**

FPU Busy bit B=0 FPU free

FPU Busy bit B=1 FPU busy.





# X87 Tag register



Each data register of x87 FPU has an associated 2 bit tag contained in a tag register.

The code in the tag register indicates the type of the value of the corresponding data register.

Tag is associated with data register

00- R has a valid value

01-R has a zero

10-R has a special value (Nan,+/- infinity, renormalized number)

11-R is empty.

Program to Illustrate Status Register, Tag register and Control register.

```
1 .section .bss
2
3
4 .lcomm status,2
5
6 .section .text
7
8 .globl _start
9 _start : nop
10
11     fstsw status
12
13     movl $1,%eax
14     movl $0,%ebx
15     int $0x80
16
```

```
student@student-OptiPlex-3020: ~
0x0000000000000000000000000000000000000000000000000000000000000000, 0x00000000
ymm2 {v8_float = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0}, v4_double = {0x0, 0x0, 0x0, 0x0}, v32_int8 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0}, v8_int32 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0}}
---Type <return> to continue, or q <return> to quit
[1]+  Stopped                  gdb -q fs
student@student-OptiPlex-3020:~$ gdb -q fs
Reading symbols from fs...done.
(gdb) break *_start+1
Breakpoint 1 at 0x8048075: file fs.s, line 11.
(gdb) run
Starting program: /home/student/fs

Breakpoint 1, _start () at fs.s:11
11         fstsw status
(gdb) s
13         movl $1,%eax
(gdb) print/x $ftag
$1 = 0xffff
(gdb) print/x $fctrl
$2 = 0x37f
(gdb) print/x $fstat
$3 = 0x0
(gdb)
```

## **Fadd**

Add two floating point number available in stack registers

Example `fadd %st(1),%st*0)`

Here `%st(0)` is added with `%st(1)` and result stored in **destination** stack `%st(0)`

## **Fiadd**

It is used to add add an integer stored in memory location to the floating point number on the top of stack.

SJMurchite\_SYB\_Unit5 syllabus

## **Faddr**

Add two floating point number available in stack registers

Example `faddr %st(1),%st(0)`

Here `%st(0)` is added with `%st(1)` and result stored in **source** stack register `%st(0)`

## **Faddp**

Add two floating point number available in stack registers and pop the top of stack.

Example `faddp %st(1),%st(0)`

here `%st(0)` is added with `%st(1)` and remove stack registers and put result in top of stack.

```
.section .data
value1:
    .float 43.65
value2:
    .int 22
value3:
    .float 76.34
value4:
    .float 3.1
value5:
    .float 12.43
value6:
    .int 6
value7:
    .float 140.2
value8:
    .float 94.21
.section .text
```

```

_start:

```

nop

finit

flds value1

fidiv value2

flds value3

flds value4

```
fmul %st(1), %st(0)
```

```
fadd %st(2), %st(0)
```

flds value5

fimul value6

flds value7

flds value8

fdivrp

```
fsubr %st(1), %st(0)
```

```
fdivr %st(2), %st(0)
```

```
movl $1,%eax
```

```
movl $0,%ebx
```

```
int $0x80
```

[illegible]

SJMurchite\_SYB\_Unit5\_syllabus

6
238.63809

13
3.264907

# Advanced Floating-Point Maths

- `.section .data`
- `value1:`
- `.float 395.21`
- `value2:`
- `.float -9145.290`
- `value3:`
- `.float 64.0`
- `.section .text`
- `.globl _start`
- `_start:`
- `nop`
- `finit`
- `flds value1`
- `fchs`
- `flds value2`
- `fabs`
- `flds value3`
- `fsqrt`
- `movl $1, %eax`
- `movl $0, %ebx`
- `int $0x80`

**Fchs** this instruction is used to change the sign of the input argument on the top of stack i.e `st(0)`

**Fabs** this instruction is used to compute the absolute value of register `st(0)`

**Fsqrt** this instruction is used to compute square root of a number on top of stack `st(0)` and returns the result in register `st(0)` overwriting the value stored previously

$$\frac{x_1}{x_2} = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

$$x^2 - 4x + 3$$

Roots of a quadratic equation

```
.section .text

.globl _start
_start : nop

    flds a
    fimul val1
    flds c
    fmul %st(1),%st(0)
    flds b
    fmul %st(0),%st(0)
    fsubp %st(1),%st(0)
    fsqrt
    flds b
    fchs
    fadd %st(1),%st(0)
    flds b
    fchs
    fsub %st(2),%st(0)
    flds a
    fimul val2
    fsts val5
    fdivr %st(2),%st(0)
    fsts val3
    flds val5
    fdivr %st(2),%st(0)
    fsts val4

    movl $1,%eax
    movl $0,%ebx
    int $0x80
```

```
(gdb) s
38          flds a
(gdb) s
39          fimul val2
(gdb) s
40          fsts val5
(gdb) s
41          fdivr %st(2),%st(0)
(gdb) s
42          fsts val3
(gdb) s
43          flds val5
(gdb) x/f &val3
0x80490f8 <val3>:          3
(gdb) s
44          fdivr %st(2),%st(0)
(gdb) s
45          fsts val4
(gdb) s
47          movl $1,%eax
(gdb) x/f &val4
0x80490fc <val4>:          1
(gdb) █
```

constant loading x87 FPU instructions

in the x87 instruction set, there are certain instructions that can load commonly used constants on the stack top. These instructions are the following

Fld1	$st(0)=1.0$
Fldz	$st(0)=0.0$
Fldpi	$st(0)=3.14$
Fldl2e	$\log_2 e$
Fldln2	$\log_e 2$
Fldl2t	$\log_2 10$
Fldlg2	$\log_{10} 2$



## Trigonometric x87 FPU instructions

In the x87 FPU instruction set, following instructions are available and can compute trigonometric functions for arguments stored on the stack top

- 
- **Fsin** It compute sin of angle in radians provided in st(0)
- 
- **Fcos** It compute cosine of angle in radians provided in st(0)
- 
- **Fsincos** It compute and leaves two values in the register stack.
- 
- **Fptan** It removes the angle provided on top of the register stack and computers its tangent.

$$\text{radians} = (\text{degrees} * \pi) / 180$$

- .section .data
- degree1:
- .float 90.0
- val180:
- .int 180
- .section .bss
- .lcomm radian1, 4
- .lcomm result1, 4
- .lcomm result2, 4
- .section .text
- .globl \_start
- \_start:
- nop
- finit
- flds degree1
- fidivs val180
- fldpi
- fmul %st(1), %st(0)
- fsts radian1
- fsin
- fsts result1
- flds radian1
- fcos
- fsts result2
- movl \$1, %eax
- movl \$0, %ebx
- int \$0x80

## logarithmic x87 FPU instructions

there are two instructions in x87 FPU instruction set to compute logarithm

**Fyl2x**

it compute  $y \log_2 x$  with  $y$  and  $x$  being in register  $st(1)$  and  $st(0)$  respectively.

SJMurchite\_SYB\_Unit5 syllabus

**fyl2xpl**

it compute  $y \log_2 (x + 1)$  with  $y$  and  $x$  being in register  $st(1)$  and  $st(0)$  respectively

Both the instructions take two implied arguments  $x$  and  $y$  on the register stack in registers  $st(0)$  and  $st(1)$  respectively.

After successful execution of this instruction ,both operands are removed and result is pushed on the register stack.

$$\log_b X = (1/\log_2 b) * \log_2 X$$

- .section .data
- value:
- .float 12.0
- base:
- .float 10.0
- .section .bss
- .lcomm result, 4
- .section .text
- .globl \_start
- \_start:
- nop
- finit
- fldl
- flds base
- fyl2x
- fldl
- fdivp
- flds value
- fyl2x
- fstps result
- movl \$1, %eax
- movl \$0, %ebx
- int \$0x80

# SIMD Technology

- **Introduction**

- Intel introduced multimedia extension (MMX) instruction set with Pentium processors.
- Instruction operate simultaneously on multiple data values.
- Application include image processing, voice and data communication.
- In Pentium III, intel introduced streaming SIMD extension (SSE) instruction set.
- While MMX instructions operate on integer data, SSE instructions operate on floating point data.
- The SSE instruction set is targeted at applications that operate on large arrays of floating point numbers such as 3D graphics, video encoding and decoding etc.

# SIMD Environment

- SIMD technology provides additional way to define integers.
- It Perform arithmetic operations on a group of multiple integers simultaneously
- SIMD architecture uses packed data type
- A packed integer is series of bytes that can represent more than one integer.
- SIMD instruction set provide a few register called MMX registers and XMM registers.
- MMX registers includes 8,64 bit registers named mm0-mm7.
- Data register are 80 bit wide, out of which only 64 bits are used by instructions in the SIMD instruction set.
- It is recommended to use finit instruction to initialize x87 FPU at the time of switching from SIMD to x87 environment.
- XMM registers includes 8,128 bit registers named xmm0-xmm7.
- XMM registers support integer and floating point data types.
- Supported floating point data type include 4 packed single precision floating point numbers.

## Loading and retrieving packed integer values

```
.section .data
```

```
packedvalue1:
```

```
.byte 10, 20, -30, 40, 50, 60, -70, 80
```

```
packedvalue2:
```

```
.short 10, 20, 30, 40
```

```
packedvalue3:
```

```
.int 10, 20
```

```
.section .text
```

```
.globl _start
```

```
_start:
```

```
movq packedvalue1, %mm0
```

```
movq packedvalue2, %mm1
```

```
movq packedvalue3, %mm2
```

# MMX addition and subtraction instructions

- With normal addition and subtraction with general-purpose registers, if an overflow condition exists from the operation, the EFLAGS register is set to indicate the overflow condition
- when using MMX addition or subtraction, you must decide ahead of time what the processor should do in case of overflow conditions within the operation.
- You can choose from three overflow methods for performing the mathematical operations:
  1. Wraparound arithmetic
  2. Signed saturation arithmetic
  3. Unsigned saturation arithmetic



# Examples of SIMD byte operations with saturations

A	B	Operation	Wrap around	Signed saturation	Unsigned saturation
0xA3	0xC2	A+B	0x65	0x80	0xFF
0x57	0xB2	A+B	0x09	0x09	0xFF
0x78	0x40	A+B	0xB8	0x7F	0xB8
0x40	0x72	A-B	0xCE	0xCE	0x00

**Wrap around-** Whenever the result is larger than 0xFF, Extra bit is dropped. Result is treated as module 256.

**Signed saturation-** when -93 and -62 in decimal added. The summation of two numbers will be -155 which is below the minimum possible value 0x80

**Unsigned saturation-** When 163 and 194 in decimal added the result is 357, which again cannot represent in 8 bit unsigned number format. So the result is 0xFf

int \$0x80

value1:

value2:

```
.section .bss
```

.lcomm result, 8

```
.section .text
```

## .globl \_start

```
_start: nop
```

```
movq value2, %mm1
```

```
movq %mm0, result
```

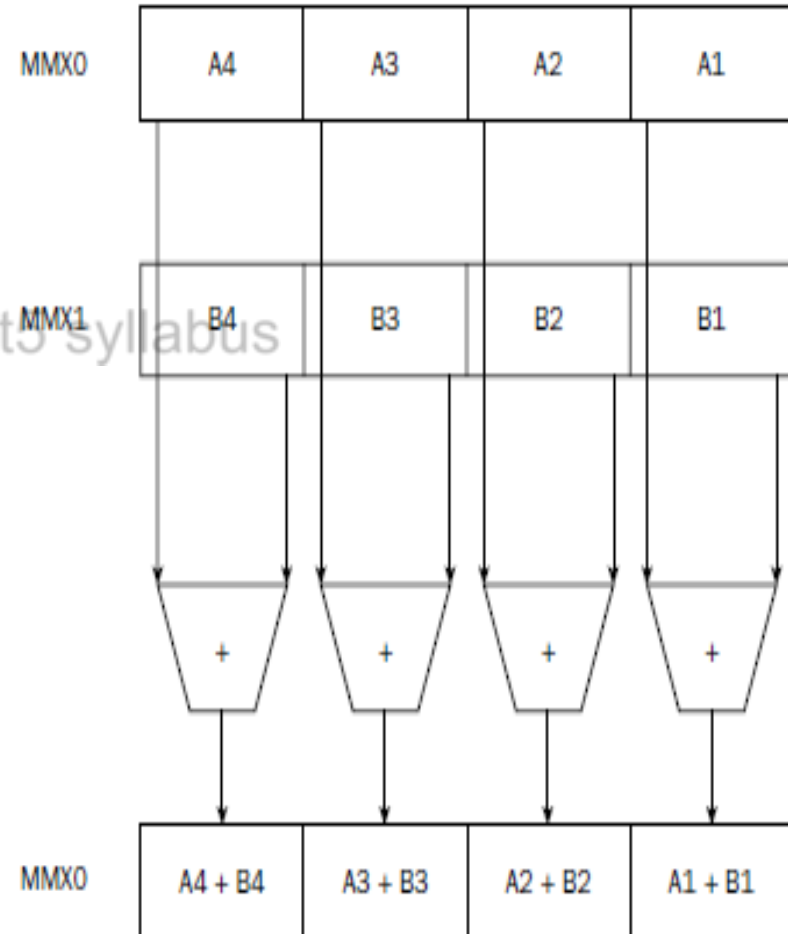
```
movl $1, %eax
```

```
movl $0, %ebx
```

int \$0x80

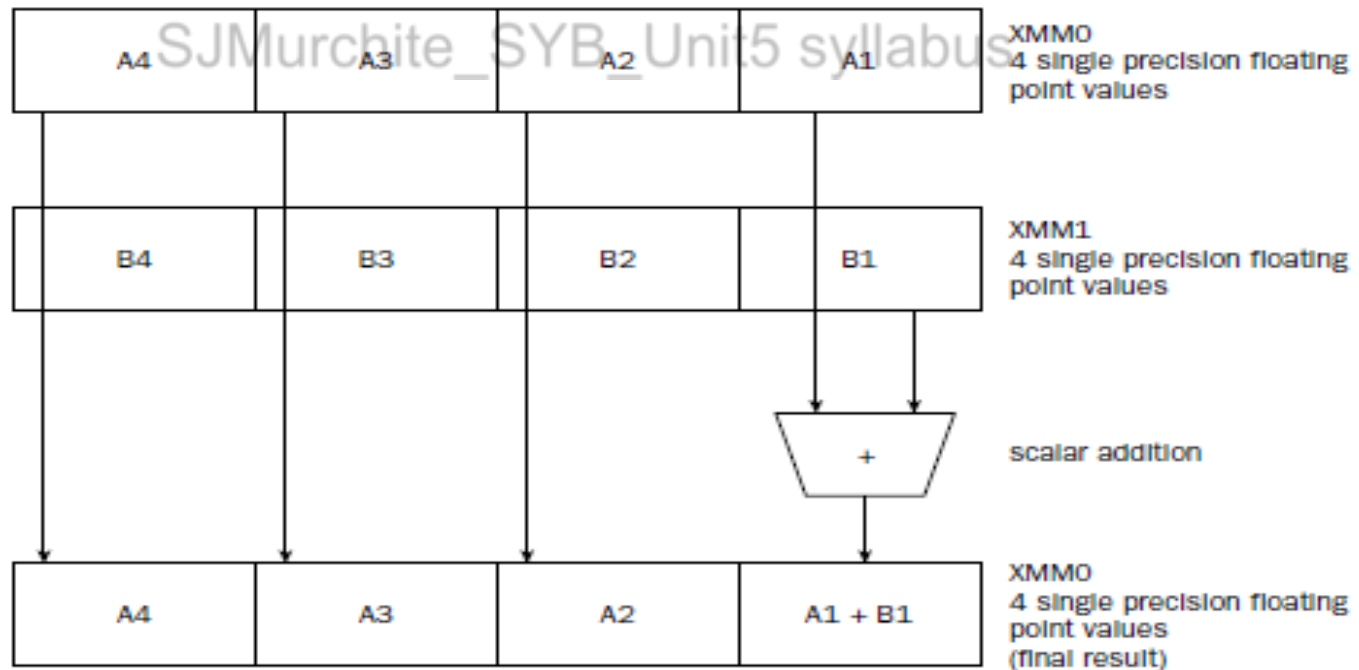
## PADDSD.....with sign saturation

## PADDUSD.....with unsigned saturation



# SSE Instructions

- Main purpose of the SSE technology is to perform SIMD operations on floating point data.



# DATA TRANSFER IN SSE TECHNOLOGY

- `.section .data`
- `.align 16`
- `value1:`
- `.float 12.34, 2345.543, -3493.2, 0.4491`
- `.section .text`
- `.globl _start`
- `_start:`
- `movaps value1, %xmm0`

`.align` directive instructs the gas assembler to align the data on a specific memory boundary.  
It takes a single operand, the size of the memory boundary on which to align the data

# Addition Example of using SSE arithmetic instructions

- `.section .data`
- `.align 16`
- `value1:`
- `.float 12.34, 2345., -93.2, 10.44`
- `value2:`
- `.float 39.234, 21.4, 100.94, 10.56`
- `.section .bss`
- `.lcomm result, 16`
- `.section .text`
- `.globl _start`
- `_start:`
- `nop`
- `movaps value1, %xmm0`
- `movaps value2, %xmm1`
- `addps %xmm1, %xmm0`
- `movaps %xmm0, result`
- `movl $1, %eax`
- `movl $0, %ebx`
- `int $0x80`

SJMurchite\_SYB\_Unit5 syllabus

SJMurchite\_SYB\_Unit5 syllabus

# Thank you